


Unsupervised learning

Clustering

DSE 210

Clustering in \mathbb{R}^d



Two common uses of clustering:

- **Vector quantization**

Find a finite set of representatives that provides good coverage of a complex, possibly infinite, high-dimensional space.

- **Finding meaningful structure in data**

Finding salient grouping in data.

Widely-used clustering methods

- ① K -means and its many variants
- ② EM for mixtures of Gaussians
- ③ Agglomerative hierarchical clustering


The k -means optimization problem

Vector quantization

- Input: Points $x_1, \dots, x_n \in \mathbb{R}^d$; integer k
- Output: "Centers", or representatives, $\mu_1, \dots, \mu_k \in \mathbb{R}^d$
- Goal: Minimize average squared distance between points and their nearest representatives:

$$\text{cost}(\mu_1, \dots, \mu_k) = \sum_{i=1}^n \min_j \|x_i - \mu_j\|^2$$

↑ squared distance
from x_i to its closest rep.




The k -means optimization problem

- Input: Points $x_1, \dots, x_n \in \mathbb{R}^d$; integer k
- Output: “Centers”, or representatives, $\mu_1, \dots, \mu_k \in \mathbb{R}^d$
- Goal: Minimize average squared distance between points and their nearest representatives:

$$\text{cost}(\mu_1, \dots, \mu_k) = \sum_{i=1}^n \min_j \|x_i - \mu_j\|^2$$


$k=3$



The k -means optimization problem

- Input: Points $x_1, \dots, x_n \in \mathbb{R}^d$; integer k
- Output: “Centers”, or representatives, $\mu_1, \dots, \mu_k \in \mathbb{R}^d$
- Goal: Minimize average squared distance between points and their nearest representatives:


$$\text{cost}(\mu_1, \dots, \mu_k) = \sum_{i=1}^n \min_j \|x_i - \mu_j\|^2$$



The k -means optimization problem

- Input: Points $x_1, \dots, x_n \in \mathbb{R}^d$; integer k
- Output: “Centers”, or representatives, $\mu_1, \dots, \mu_k \in \mathbb{R}^d$
- Goal: Minimize average squared distance between points and their nearest representatives:

$$\text{cost}(\mu_1, \dots, \mu_k) = \sum_{i=1}^n \min_j \|x_i - \mu_j\|^2$$




Centers carve \mathbb{R}^d into k **convex** regions: μ_j 's region consists of points for which it is the closest center.

Lloyd's k -means algorithm

NP-hard optimization problem. Heuristic: “ k -means algorithm”.


- Initialize centers μ_1, \dots, μ_k in some manner.
- Repeat until convergence:
 - Assign each point to its closest center.
 - Update each μ_j to the mean of the points assigned to it.



Lloyd's k -means algorithm

NP-hard optimization problem. Heuristic: “ k -means algorithm”.


- Initialize centers μ_1, \dots, μ_k in some manner.
- Repeat until convergence:
 - Assign each point to its closest center.
 - Update each μ_j to the mean of the points assigned to it.



Lloyd's k -means algorithm

NP-hard optimization problem. Heuristic: “ k -means algorithm”.


- Initialize centers μ_1, \dots, μ_k in some manner.
- Repeat until convergence:
 - Assign each point to its closest center.
 - Update each μ_j to the mean of the points assigned to it.



Lloyd's k -means algorithm

NP-hard optimization problem. Heuristic: “ k -means algorithm”.


- Initialize centers μ_1, \dots, μ_k in some manner.
- Repeat until convergence:
 - Assign each point to its closest center.
 - Update each μ_j to the mean of the points assigned to it.



Lloyd's k -means algorithm

NP-hard optimization problem. Heuristic: “ k -means algorithm”.


- Initialize centers μ_1, \dots, μ_k in some manner.
- Repeat until convergence:
 - Assign each point to its closest center.
 - Update each μ_j to the mean of the points assigned to it.



Lloyd's k -means algorithm

NP-hard optimization problem. Heuristic: “ k -means algorithm”.




- Initialize centers μ_1, \dots, μ_k in some manner.
- Repeat until convergence:
 - Assign each point to its closest center.
 - Update each μ_j to the mean of the points assigned to it.



Lloyd's k -means algorithm

NP-hard optimization problem. Heuristic: “ k -means algorithm”.



- Initialize centers μ_1, \dots, μ_k in some manner.
- Repeat until convergence:
 - Assign each point to its closest center.
 - Update each μ_j to the mean of the points assigned to it.



Each iteration reduces the cost \Rightarrow convergence to a local optimum.

Initialization matters

Global optimum: the solution (set of "centers" μ_1, \dots, μ_k) with the smallest achievable cost.



Local optimum but
not global optimum

Initializing the k -means algorithm

Typical practice: choose k data points at random as the initial centers.

Initializing the k -means algorithm

Typical practice: choose k data points at random as the initial centers.

Another common trick: start with extra centers, then prune later.

Initializing the k -means algorithm

Typical practice: choose k data points at random as the initial centers.

Another common trick: start with extra centers, then prune later.


A particularly good initializer: k -means++ (2006)

- Pick a data point x at random as the first center
- Let $C = \{x\}$ (centers chosen so far)
- Repeat until desired number of centers is attained:
 - Pick a data point x at random from the following distribution:

$$\Pr(x) \propto \text{dist}(x, C)^2,$$

where $\text{dist}(x, C) = \min_{z \in C} \|x - z\|$

- Add x to C



Two common uses of clustering

- Vector quantization

Find a finite set of representatives that provides good coverage of a complex, possibly infinite, high-dimensional space.




- Finding meaningful structure in data

Finding salient grouping in data.


Representing images using k -means codewords

Vector quantization
example

How to represent a collection of images as fixed-length vectors?



Each patch becomes a vector:



apply k -means to the entire collection

→ yields K "representative" patches
(eg. $K=100$)

"Bag of patches"

Representing images using k -means codewords

How to represent a collection of images as fixed-length vectors?



- Take all $\ell \times \ell$ patches in all images. Extract features for each.
- Run k -means on this entire collection to get k centers.
- Now associate any image patch with its nearest center.
- Represent an image by a histogram over $\{1, 2, \dots, k\}$.

Looking for natural groups in data

“Animals with attributes” data set

- 50 animals: antelope, grizzly bear, beaver, dalmatian, tiger, ...
- 85 attributes: longneck, tail, walks, swims, nocturnal, forager, desert, bush, plains, ...
- Each animal gets a score (0 – 100) along each attribute
- 50 data points in \mathbb{R}^{85}

Apply k -means with $k = 10$ and look at grouping obtained.

Two different solutions starting from random initializations

- 1 zebra
- 2 spider monkey, gorilla, chimpanzee
- 3 tiger, leopard, wolf, bobcat, lion
- 4 hippopotamus, elephant, rhinoceros
- 5 killer whale, blue whale, humpback whale, seal, walrus, dolphin
- 6 giant panda
- 7 skunk, mole, hamster, squirrel, rabbit, bat, rat, weasel, mouse, raccoon
- 8 antelope, horse, moose, ox, sheep, giraffe, buffalo, deer, pig, cow
- 9 beaver, otter
- 10 grizzly bear, dalmatian, persian cat, german shepherd, siamese cat, fox, chihuahua, polar bear, collie

- 1 zebra
- 2 spider monkey, gorilla, chimpanzee
- 3 tiger, leopard, fox, wolf, bobcat, lion
- 4 hippopotamus, elephant, rhinoceros, buffalo, pig
- 5 killer whale, blue whale, humpback whale, seal, otter, walrus, dolphin
- 6 dalmatian, persian cat, german shepherd, siamese cat, chihuahua, giant panda, collie
- 7 beaver, skunk, mole, squirrel, bat, rat, weasel, mouse, raccoon
- 8 antelope, horse, moose, ox, sheep, giraffe, deer, cow
- 9 hamster, rabbit
- 10 grizzly bear, polar bear

Streaming and online computation

Streaming computation: for data sets too large to fit in memory.

- Make one pass (or maybe a few passes) through the data.
- On each pass:
 - See data points one at a time, in order.
 - Update models/parameters along the way.
- There is only enough space to store a tiny fraction of the data, or a perhaps short summary.

Streaming and online computation

Streaming computation: for data sets too large to fit in memory.

- Make one pass (or maybe a few passes) through the data.
- On each pass:
 - See data points one at a time, in order.
 - Update models/parameters along the way.
- There is only enough space to store a tiny fraction of the data, or a perhaps short summary.

Online computation: an even more lightweight setup, for data that is continuously being collected.

- Initialize a model.
- Repeat forever:
 - See a new data point.
 - Update model if need be.

Example: sequential k -means

- ① Set the centers μ_1, \dots, μ_k to the first k data points
- ② Set their counts to $n_1 = n_2 = \dots = n_k = 1$
- ③ Repeat, possibly forever:
 - Get next data point x
 - Let μ_j be the center closest to x
 - Update μ_j and n_j :


$$\mu_j = \frac{n_j \mu_j + x}{n_j + 1} \quad \text{and} \quad n_j = n_j + 1$$

average x
into μ_j

K-means: the good and the bad

The good:

- Fast and easy.
- Effective in quantization.




The bad:


- Geared towards data in which the clusters are spherical, and of roughly the same radius.

Is there is a similarly-simple algorithm in which clusters of more general shape are accommodated?

Mixtures of Gaussians




Mixtures of Gaussians



Mixtures of Gaussians

K-means: each cluster only gets a center
Now: each cluster gets its own Gaussian distribution



Each of the k clusters is specified by:

- a Gaussian distribution $P_j = N(\mu_j, \Sigma_j)$
- a mixing weight π_j

Overall distribution over \mathbb{R}^d : a **mixture of Gaussians**

$$\Pr(x) = \pi_1 P_1(x) + \cdots + \pi_k P_k(x)$$

The clustering task

We are given data $x_1, \dots, x_n \in \mathbb{R}^d$.

For any mixture model π_1, \dots, π_k , $P_1 = N(\mu_1, \Sigma_1), \dots, P_k = N(\mu_k, \Sigma_k)$,

$$\begin{aligned} & \Pr(\text{data} \mid \pi_1 P_1 + \dots + \pi_k P_k) \\ &= \prod_{i=1}^n (\underbrace{\pi_1 P_1(x_i) + \dots + \pi_k P_k(x_i)}_{\text{probability of } x_i \text{ under the mixture model}}) \\ &= \prod_{i=1}^n \left(\sum_{j=1}^k \frac{\pi_j}{(2\pi)^{d/2} |\Sigma_j|^{1/2}} \exp \left(-\frac{1}{2} (x_i - \mu_j)^T \Sigma_j^{-1} (x_i - \mu_j) \right) \right) \end{aligned}$$

Find the **maximum-likelihood mixture of Gaussians**: parameters $\{\pi_j, \mu_j, \Sigma_j : j = 1 \dots k\}$ maximizing this function.

complicated optimization problem

Optimization surface

Minimize the negative log-likelihood,

$$\sum_{i=1}^n \ln \left(\sum_{j=1}^k \frac{\pi_j}{(2\pi)^{d/2} |\Sigma_j|^{1/2}} \exp \left(-\frac{1}{2} (x_i - \mu_j)^T \Sigma_j^{-1} (x_i - \mu_j) \right) \right)$$

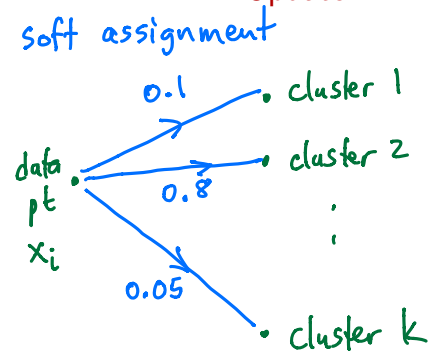


The EM algorithm

- 1 Initialize π_1, \dots, π_k and $P_1 = N(\mu_1, \Sigma_1), \dots, P_k = N(\mu_k, \Sigma_k)$.
- 2 Repeat until convergence:
 - Assign each point x_i fractionally between the k clusters:

$$w_{ij} = \Pr(\text{cluster } j \mid x_i) = \frac{\pi_j P_j(x_i)}{\sum_{\ell} \pi_{\ell} P_{\ell}(x_i)}$$

- Update mixing weights, means, and covariances:



$$\pi_j = \frac{1}{n} \sum_{i=1}^n w_{ij}$$

$$\mu_j = \frac{1}{n\pi_j} \sum_{i=1}^n w_{ij} x_i$$

$$\Sigma_j = \frac{1}{n\pi_j} \sum_{i=1}^n w_{ij} (x_i - \mu_j)(x_i - \mu_j)^T$$

Example


Data with 3 clusters, each with 100 points.



The data set

Example


Data with 3 clusters, each with 100 points.



k -means solution 1

Example


Data with 3 clusters, each with 100 points.



k -means solution 2

Example


Data with 3 clusters, each with 100 points.



k -means solution 3

Example


Data with 3 clusters, each with 100 points.



EM for mixture of Gaussians

Hierarchical clustering


Choosing the number of clusters (k) is difficult.





Often: no single right answer, because of multiscale structure.

Hierarchical clustering avoids these problems.

Example: gene expression data



The single linkage algorithm



- Start with each point in its own, singleton, cluster
- Repeat until there is just one cluster:
 - Merge the two clusters with the closest pair of points
- Disregard singleton clusters


Linkage methods

- Start with each point in its own, singleton, cluster
- Repeat until there is just one cluster:
 - Merge the two “closest” clusters

Linkage methods

- Start with each point in its own, singleton, cluster
- Repeat until there is just one cluster:
 - Merge the two “closest” clusters

How to measure distance between two clusters C and C' ?



- Single linkage

$$\text{dist}(C, C') = \min_{x \in C, x' \in C'} \|x - x'\|$$

- Complete linkage

$$\text{dist}(C, C') = \max_{x \in C, x' \in C'} \|x - x'\|$$

Average linkage

Three commonly-used variants:

- ① Average pairwise distance between points in the two clusters

$$\text{dist}(C, C') = \frac{1}{|C| \cdot |C'|} \sum_{x \in C} \sum_{x' \in C'} \|x - x'\|$$

- ② Distance between cluster centers

$$\text{dist}(C, C') = \|\text{mean}(C) - \text{mean}(C')\|$$

- ③ Ward's method: the increase in k -means cost occasioned by merging the two clusters

$$\text{dist}(C, C') = \frac{|C| \cdot |C'|}{|C| + |C'|} \|\text{mean}(C) - \text{mean}(C')\|^2$$